

LLM Provider Ecosystem & Model Releases

Comprehensive analysis of prominent models released within the last 12 months

Data compiled as of July 2026 • Market Intelligence Analysis

MODEL PROVIDER	MODEL	RELEASE DATE	CUT-OFF DATE	COST OF TOKENS (PER 1M)	CONTEXT WINDOW	TYPE	WHEN TO USE THIS MODEL
OpenAI	GPT-5.5	April 24, 2026	Dec 1, 2025	\$5.00 Input \$30.00 Output	1,000,000	Proprietary Dense	Complex multi-step reasoning, advanced agentic orchestration, and high-tier coding.
Anthropic	Claude Sonnet 5	June 30, 2026	Mid 2025	~\$3.00 Input ~\$15.00 Output	200,000	Proprietary Multimodal	Long-form content parsing, multi-layered enterprise workflows, and deep codebase generation.
Google	Gemini 3.5 Flash	May 19, 2026	Late 2025	~\$0.075 Input ~\$0.30 Output	1,000,000	Proprietary MoE	High-frequency, high-volume tasks requiring ultra-low latency and vast context lengths (video/audio).
DeepSeek	DeepSeek V4	April 24, 2026	Late 2025	\$0.14 Input \$0.28 Output	1,000,000	Open-Weights MoE	High-performance reasoning and code generation on an aggressive development budget.
Microsoft	Phi-4-reasoning-vision-15B	March 2026	Late 2025	Free (Open Weights)	128,000	Open-Weights SLM	Local deployment, edge devices, or resource-constrained environments requiring vision-language reasoning.
Mistral AI	Mistral Large 3	December 2, 2025	Mid 2025	~\$2.00 Input ~\$6.00 Output	256,000	Open-Weights MoE	Sovereign enterprise setups requiring native multilingual processing (40+ languages) and document analysis.

MODEL PROVIDER	MODEL	RELEASE DATE	CUT-OFF DATE	COST OF TOKENS (PER 1M)	CONTEXT WINDOW	TYPE	WHEN TO USE THIS MODEL
Cohere	Command A+	May 20, 2026	Late 2025	Free (Apache 2.0)	128,000	Open-Weights MoE	Advanced enterprise Retrieval-Augmented Generation (RAG) that relies heavily on native, grounded citations.
xAI	Grok 4.5	June 28, 2026	Early 2026	Tiered API (Subscription)	512,000	Proprietary Large	Scenarios requiring real-time cross-referenced network information and intensive mathematical reasoning.
Alibaba Cloud	Qwen3.7-Max	May 20, 2026	Early 2026	\$2.50 Input \$7.50 Output	128,000	Proprietary MoE	Repository-level problem solving, complex agentic coding, and strong multi-lingual support across Asian languages.
Meta	Muse Spark Image	July 7, 2026	Mid 2025	Free via Meta AI (Ad Suite)	N/A (Image Modality)	Proprietary Vision	Complex creative image rendering, asset creation, and graphic ad generation using advanced layout instructions.

Architectural & Pricing Note:

For open-weight models (such as *Phi-4*, *Command A+*, and *Mistral Large 3*), token costs are technically free when deployed on self-hosted architecture. Commercial cloud endpoints (e.g., Hugging Face, Azure, OpenRouter) host these open weights as managed services and charge varying competitive baseline rates per million tokens.